

Data Flow Diagram (DFD)

What it shows

How raw sensor data turns into an alert or a log. One path in. One decision point. Two outputs.

The picture

```
Error parsing Mermaid diagram!  
  
Cannot read properties of null (reading 'getBoundingClientRect')
```

Steps in order

1. **Sensors → Ingest**
Read tag values. Example: Modbus/TCP or OPC-UA.
2. **Ingest → Buffer (sliding window)**
Keep the last N points per tag. Oldest drops first.
3. **Buffer → Feature Extract**
Compute simple stats per window. Mean. Std. Min. Max. Delta. Rate of change.
4. **Features → Model**
Run a lightweight model. Example: z-score. IsolationForest. One-Class SVM. Small tree. Linear.
5. **Model → Score**
Produce one number per tick. Higher means more abnormal.
6. **Score → Threshold**
Compare score to a cutoff.
7. **Decision**
 - **Yes** (score > cutoff) → **Alert**
Send message with tag name. Score. Window stats. Last values.
 - **No** → **Log**
Store score and features for learning and dashboards.

Inputs and outputs

- **Input**: time series from sensors.
- **Output A**: alert to people or tools.
- **Output B**: log row with features and score.

Knobs you can tune

- Window length N.
- Feature set.
- Model type and parameters.
- Cutoff value.

What to store

- Timestamp. Tag id. Score. Cutoff.
- Features used.
- Alert payload if fired.
- Optional raw window slice.

Metrics to watch

- Precision and recall or PR-AUC.
- Alert latency.
- False positive rate.
- Drift in feature distributions.

Edge constraints

- Run fast. Low CPU and RAM. No GPU.
- Fail closed on data gaps. If window too small then skip alert and log reason.

Failure cases and actions

- **Missing data**: hold state and log a gap event.
- **Flatline**: detect zero variance and alert as sensor fault.
- **Clock skew**: reject out-of-order points and log.