

# Explainable Artificial Intelligence

# Introduction

- AI is critical to many sectors in Information Technologies

# Introduction

- AI is critical to many sectors in Information Technologies
- Today's AI is achieving unprecedented levels of performance
- Becoming increasingly important

Better Go/Chess Player



More “Realistic” Text  
Generation



# Introduction

- AI is critical to many sectors in Information Technologies
- Today's AI is achieving unprecedented levels of performance
- Becoming increasingly important

Better Go/Chess Player



More “Realistic” Text Generation

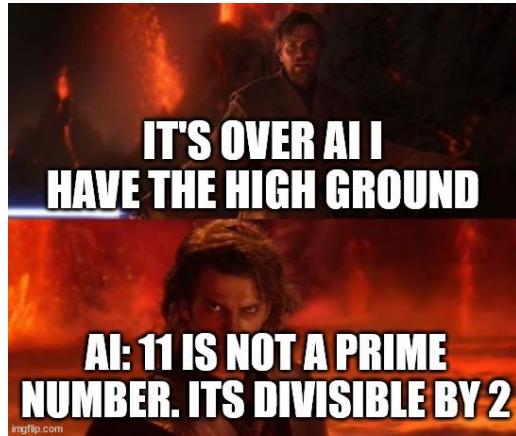


Equal to Humans at Drawing Hands



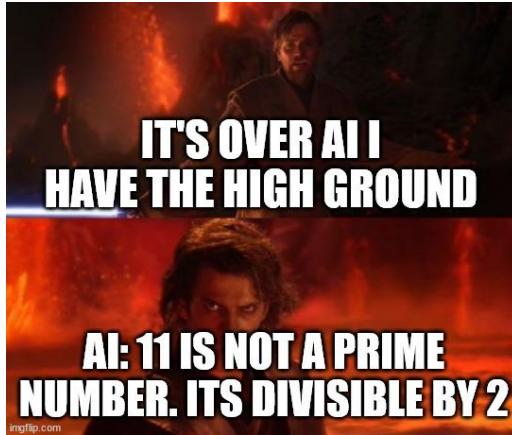
# Unfortunately

- Deployment in critical fields faces an uphill battle



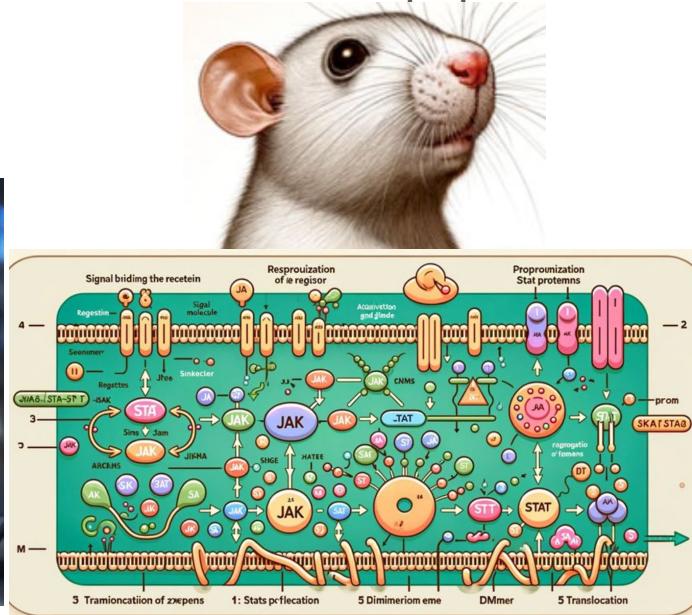
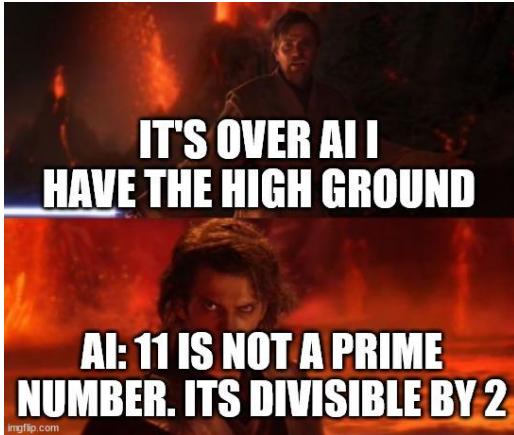
# Unfortunately

- Deployment in critical fields faces an uphill battle
- Decisions that affect humans can be difficult to accept



# Unfortunately

- Deployment in critical fields faces an uphill battle
  - Decisions that affect humans can be difficult to accept
  - Just like how “journals” shouldn’t be accepting generative ai created papers



# Unfortunately

- Deployment in critical fields faces an uphill battle
- Decisions that affect humans can be difficult to accept
- Just like how “journals” shouldn’t be accepting generative ai created papers

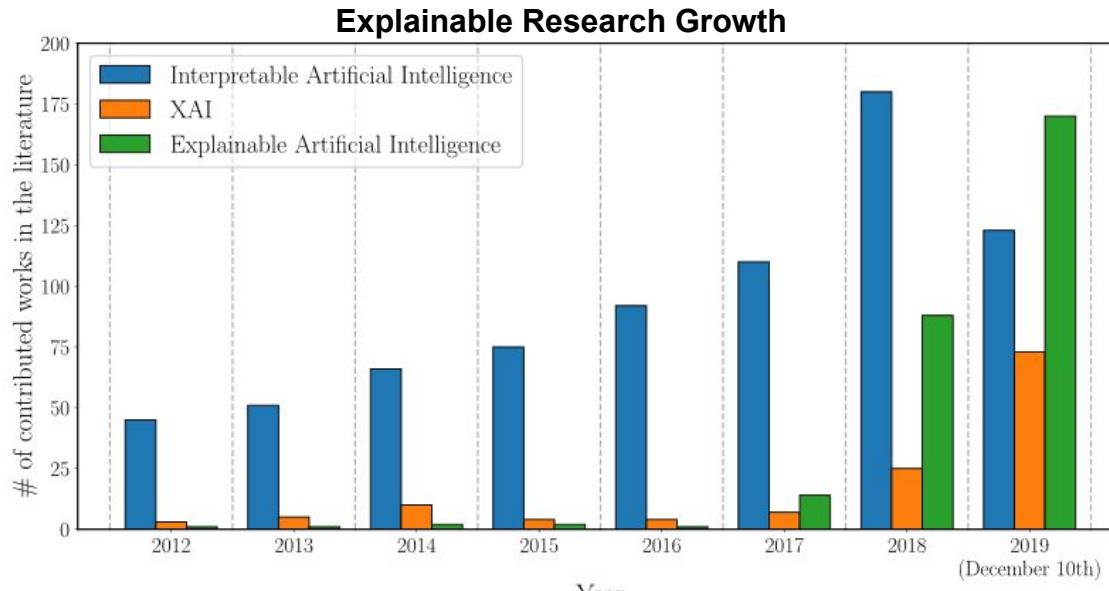
RETRACTED: Cellular functions of spermatogonial stem cells in relation to JAK/STAT signaling pathway Retracted

# Thus

- There is need in understanding decisions made by AI

# Thus

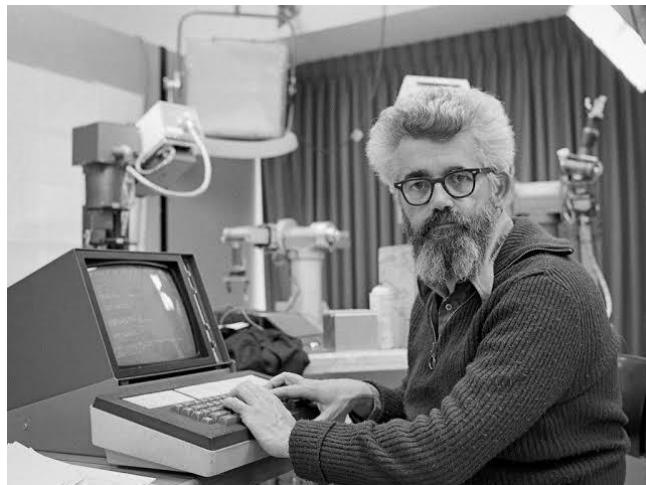
- There is need in understanding decisions made by AI
- The first AI systems were interpretable
- The current age of AI is dominated by opaque decision processes



# History Lesson

- AI dates back to the 1950's

**John McCarthy**



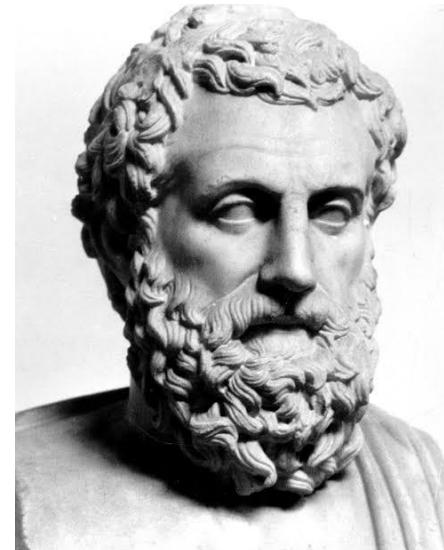
# History Lesson

- AI dates back to the 1950's
- Even further if you consider Aristotle and Greek myths

**John McCarthy**



**Aristotle**



# History Lesson

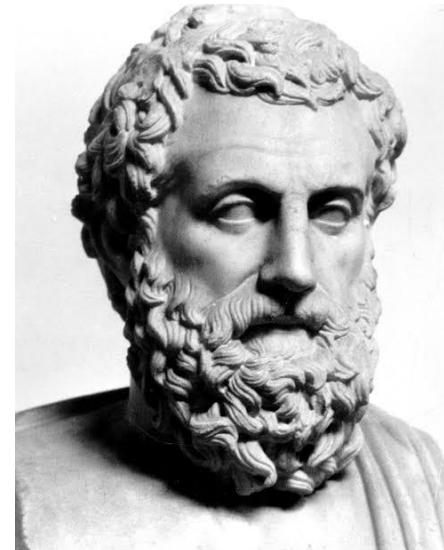
- AI dates back to the 1950's
- Even further if you consider Aristotle and Greek myths

**John McCarthy**



Old school cool

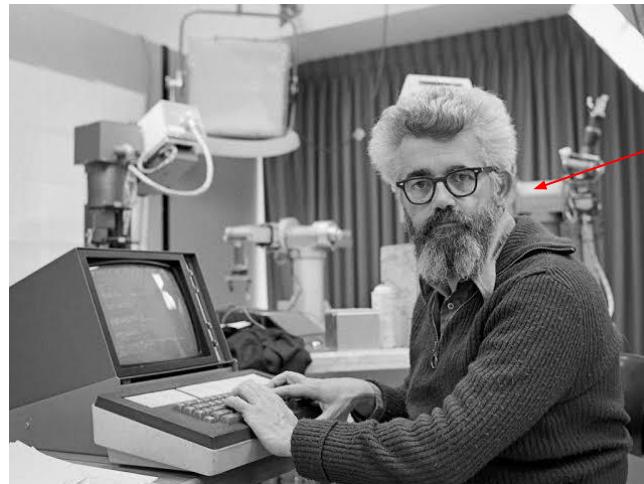
**Aristotle**



# History Lesson

- AI dates back to the 1950's
- Even further if you consider Aristotle and Greek myths

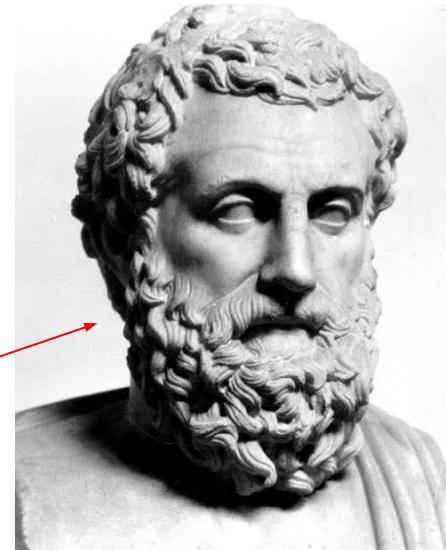
John McCarthy



Old school cool

Literally  
immortalized into  
stone

Aristotle



# History Lesson

- The “Golden Age” of AI 1956-1974
- Micro-Worlds by Marvin Minsky

*“The question of whether a computer can think is no more interesting than the question of whether a submarine can swim.”*

— Edsger Dijkstra

# History Lesson

- The “Golden Age” of AI 1956-1974
- Micro-Worlds by Marvin Minsky
  - Simplify the world into a very specific set of information

*“The question of whether a computer can think is no more interesting than the question of whether a submarine can swim.”*

— Edsger Dijkstra

# History Lesson

- The “Golden Age” of AI 1956-1974
- Micro-Worlds by Marvin Minsky
  - Simplify the world into a very specific set of information
- John McCarthy’s development of Lisp

*“The question of whether a computer can think is no more interesting than the question of whether a submarine can swim.”*

— Edsger Dijkstra

# History Lesson

- The “Golden Age” of AI 1956-1974
- Micro-Worlds by Marvin Minsky
  - Simplify the world into a very specific set of information
- John McCarthy’s development of Lisp
- The first chatbot ELIZA
- Lots of research funded by DARPA

*“The question of whether a computer can think is no more interesting than the question of whether a submarine can swim.”*

— Edsger Dijkstra

```
Welcome to
EEEEE   LL      IIII    ZZZZZZ  AAAAAA
EE      LL      II      ZZ      AA  AA
EEEEE   LL      II      ZZZ    AAAAAAAA
EE      LL      III     ZZ      AA  AA
EEEEE   LLLLLL  IIII    ZZZZZZ  AA  AA

Eliza is a mock Rogerian psychotherapist.
The original program was described by Joseph Weizenbaum in 1966.
This implementation by Norbert Landsteiner 2005.

ELIZA: Is something troubling you ?
YOU: Men are all alike.
ELIZA: What is the connection, do you suppose ?
YOU: They're always bugging us about something or other.
ELIZA: Can you think of a specific example ?
YOU: Well, my boyfriend made me come here.
ELIZA: Is it important to you that your boyfriend made you come here ?
YOU: He says I'm depressed much of the time.
ELIZA: I am sorry to hear that you are depressed.
YOU: It's true. I am unhappy.
ELIZA: Can you explain what made you unhappy ?
YOU:
```

# History Lesson

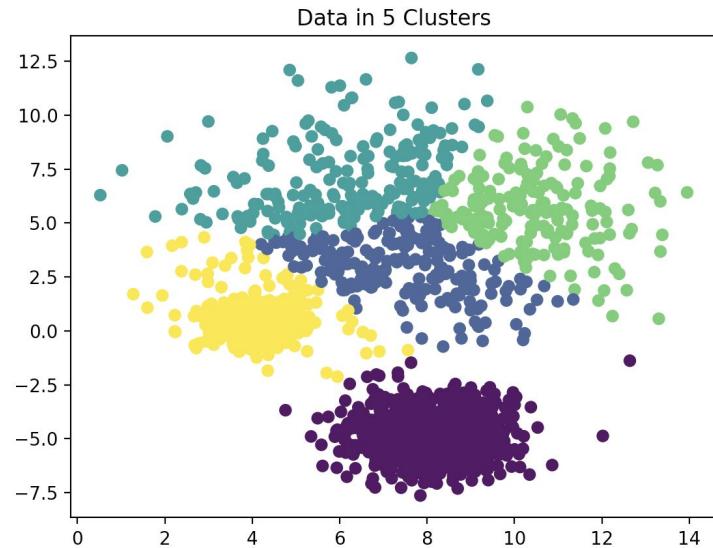
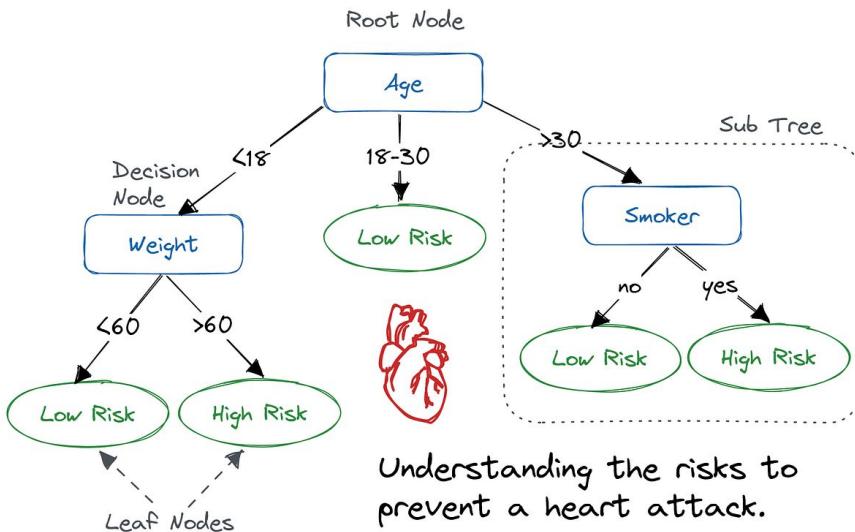
- Early AI was considered “white-box”

# History Lesson

- Early AI was considered “white-box”
- What were some examples of “white-box” AI algorithms?

# History Lesson

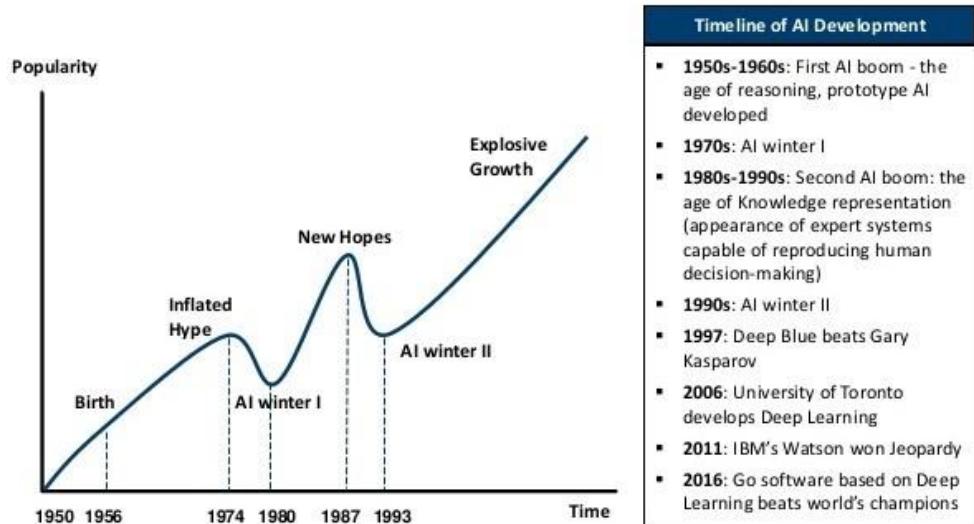
- Early AI was considered “white-box”
- Rule-based algorithms
- Linear decision
- Clustering



# History Lesson

- The first AI winter in the 1970s

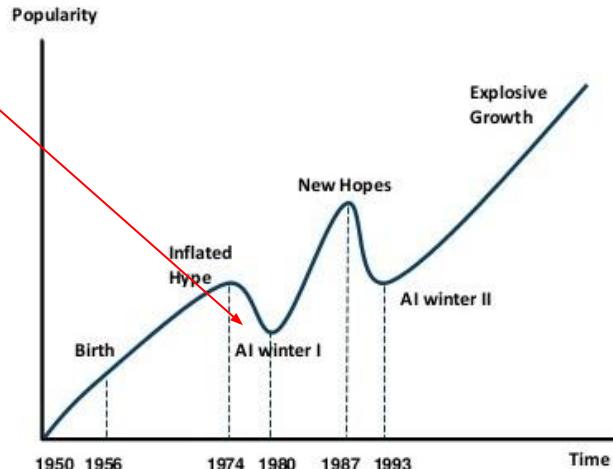
AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING”...



# History Lesson

- The first AI winter in the 1970s
- Funding dried up, popularity diminished

AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING”...



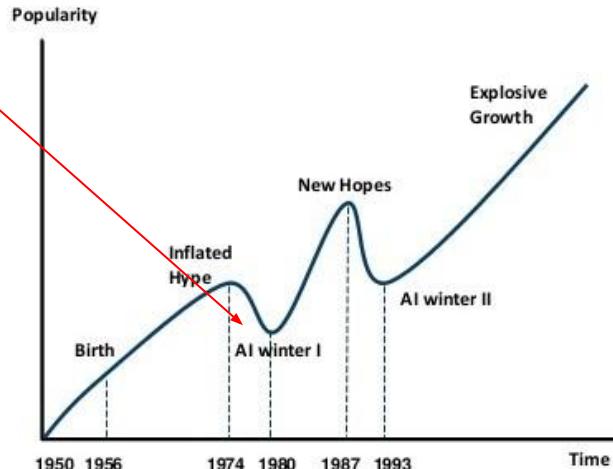
## Timeline of AI Development

- 1950s-1960s: First AI boom - the age of reasoning, prototype AI developed
- 1970s: AI winter I
- 1980s-1990s: Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making)
- 1990s: AI winter II
- 1997: Deep Blue beats Gary Kasparov
- 2006: University of Toronto develops Deep Learning
- 2011: IBM's Watson won Jeopardy
- 2016: Go software based on Deep Learning beats world's champions

# History Lesson

- The first AI winter in the 1970s
- Funding dried up, popularity diminished
- Why?

AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING”...



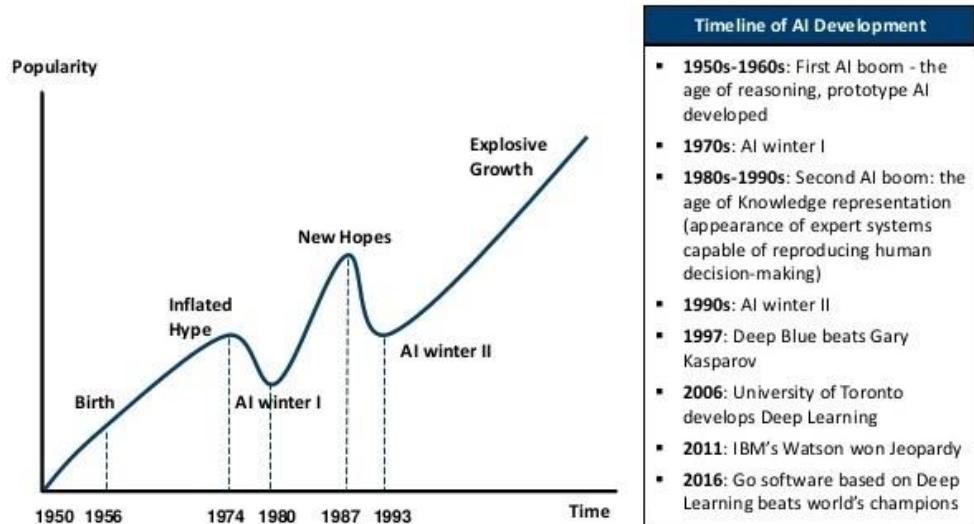
## Timeline of AI Development

- 1950s-1960s: First AI boom - the age of reasoning, prototype AI developed
- 1970s: AI winter I
- 1980s-1990s: Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making)
- 1990s: AI winter II
- 1997: Deep Blue beats Gary Kasparov
- 2006: University of Toronto develops Deep Learning
- 2011: IBM's Watson won Jeopardy
- 2016: Go software based on Deep Learning beats world's champions

# History Lesson

- The first AI winter in the 1970s
- Funding dried up, popularity diminished
- Likely
  - Technical limitations
  - Unrealistic expectations
  - No practical use at the time

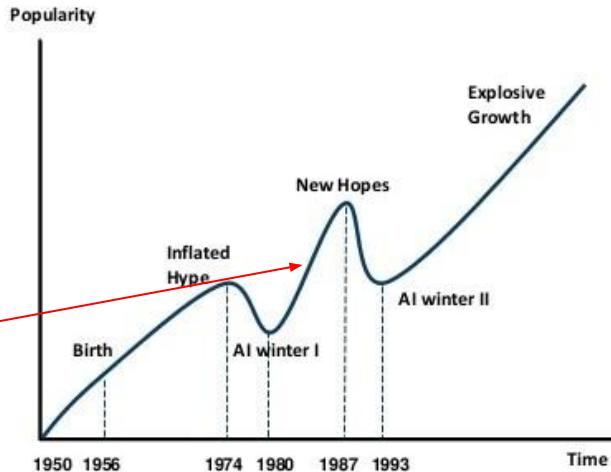
AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING”...



# History Lesson

- The first AI winter in the 1970s
- Funding dried up, popularity diminished
- Likely
  - Technical limitations
  - Unrealistic expectations
  - No practical use at the time
- Rise in expert systems

AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING”...

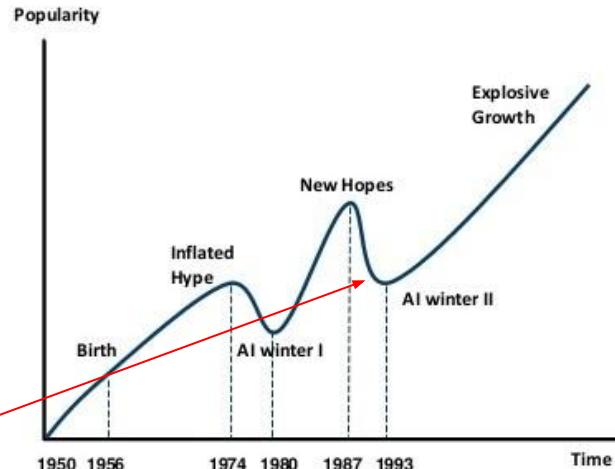


Timeline of AI Development	
▪ 1950s-1960s:	First AI boom - the age of reasoning, prototype AI developed
▪ 1970s:	AI winter I
▪ 1980s-1990s:	Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making)
▪ 1990s:	AI winter II
▪ 1997:	Deep Blue beats Gary Kasparov
▪ 2006:	University of Toronto develops Deep Learning
▪ 2011:	IBM's Watson won Jeopardy
▪ 2016:	Go software based on Deep Learning beats world's champions

# History Lesson

- The first AI winter in the 1970s
- Funding dried up, popularity diminished
- Likely
  - Technical limitations
  - Unrealistic expectations
  - No practical use at the time
- Rise in expert systems
- Fall because of technical limits

AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING”...



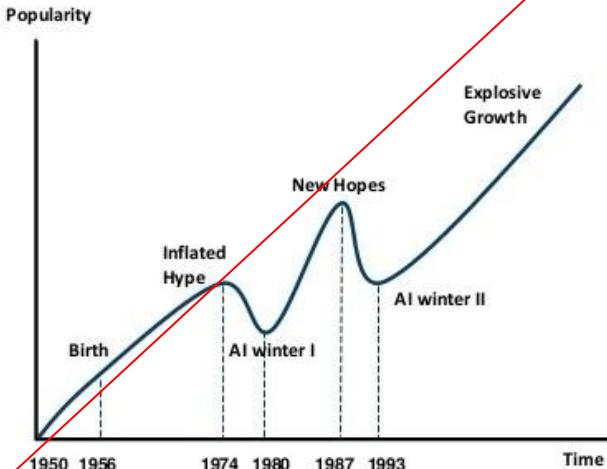
## Timeline of AI Development

- 1950s-1960s: First AI boom - the age of reasoning, prototype AI developed
- 1970s: AI winter I
- 1980s-1990s: Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making)
- 1990s: AI winter II
- 1997: Deep Blue beats Gary Kasparov
- 2006: University of Toronto develops Deep Learning
- 2011: IBM's Watson won Jeopardy
- 2016: Go software based on Deep Learning beats world's champions

# History Lesson

- The first AI winter in the 1970s
- Funding dried up, popularity diminished
- Likely
  - Technical limitations
  - Unrealistic expectations
  - No practical use at the time
- Rise in expert systems
- Fall because of technical limits
- Moore's law and Neural networks

AI HAS A LONG HISTORY OF BEING "THE NEXT BIG THING"...



Timeline of AI Development	
▪ 1950s-1960s:	First AI boom - the age of reasoning, prototype AI developed
▪ 1970s:	AI winter I
▪ 1980s-1990s:	Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making)
▪ 1990s:	AI winter II
▪ 1997:	Deep Blue beats Gary Kasparov
▪ 2006:	University of Toronto develops Deep Learning
▪ 2011:	IBM's Watson won Jeopardy
▪ 2016:	Go software based on Deep Learning beats world's champions



# Historical XAI

- Dates back to 1970's

# Historical XAI

- Dates back to 1970's
- MYCIN for medical diagnostic systems
- DENDRAL for chemical analysis
- Both used rule-based explanations
- Also "canned" XAI systems

# Historical XAI

- Dates back to 1970's
- MYCIN for medical diagnostic systems
- DENDRAL for chemical analysis
- Both used rule-based explanations
- Also "canned" XAI systems
  - Pre generated responses about domain specific decisions

# Back to Today

- Deep Neural Networks are used everywhere
- They are highly accurate
- But they are not trustworthy, tractable, or interpretable

# Back to Today

- Deep Neural Networks are used everywhere
  - They are highly accurate
  - But they are not trustworthy, tractable, or interpretable
- 
- Increase in demand for “**ethical AI**”
  - Need to understand trade off of **transparency to performance**
  - Systems can be **improved** by **understanding** them

# Back to Today

- Deep Neural Networks are used everywhere
  - They are highly accurate
  - But they are not trustworthy, tractable, or interpretable
- 
- Increase in demand for “**ethical AI**”
  - Need to understand trade off of **transparency to performance**
  - Systems can be **improved** by **understanding** them
    - This is the idea that researchers ignore/miss

# Important Terminology

- **Understandability** – characteristic of a model to make a human understand its function (how the model works) without the need for explaining its internal structure or algorithmic processes

# Important Terminology

- **Understandability** – characteristic of a model to make a human understand its function (how the model works) without the need for explaining its internal structure or algorithmic processes
- **Comprehensibility** – ability of a learned model to represent its learned knowledge to a human

# Important Terminology

- **Understandability** – characteristic of a model to make a human understand its function (how the model works) without the need for explaining its internal structure or algorithmic processes
- **Comprehensibility** – ability of a learned model to represent its learned knowledge to a human
- **Interpretability** – ability to explain or to provide meaning in understandable terms to a human

# Important Terminology

- **Understandability** – characteristic of a model to make a human understand its function (how the model works) without the need for explaining its internal structure or algorithmic processes
- **Comprehensibility** – ability of a learned model to represent its learned knowledge to a human
- **Interpretability** – ability to explain or to provide meaning in understandable terms to a human
- **Explainability** – Notion of explanations as an interface between humans and AI (Reasoning, Strengths and Weaknesses, Future behavior)

# Important Terminology

- **Understandability** – characteristic of a model to make a human understand its function (how the model works) without the need for explaining its internal structure or algorithmic processes
- **Comprehensibility** – ability of a learned model to represent its learned knowledge to a human
- **Interpretability** – ability to explain or to provide meaning in understandable terms to a human
- **Explainability** – Notion of explanations as an interface between humans and AI (Reasoning, Strengths and Weaknesses, Future behavior)
- **Transparency** – A model is considered to be transparent if by itself it is understandable

# Important Terminology

- **Understandability** – characteristic of a model to make a human understand its function (how the model works) without the need for explaining its internal structure or algorithmic processes
  - The authors believe that **Understandability** is the most crucial concept in XAI

# Important Terminology

- **Understandability** – characteristic of a model to make a human understand its function (how the model works) without the need for explaining its internal structure or algorithmic processes
  - The authors believe that **Understandability** is the most crucial concept in XAI
- **Explainability** – Notion of explanations as an interface between humans and AI (Reasoning, Strengths and Weaknesses, Future behavior)  
**Transparency** – A model is considered to be transparent if by itself it is understandable
  - I believe that the usability and transparency of an AI system is the most crucial concept in XAI
  - Understandability is also important

# Other Important Terminology

- **Local Explanation** – interpretation that provides insights into a specific instance or data point

# Other Important Terminology

- **Local Explanation** – interpretation that provides insights into a specific instance or data point
- **Global Explanation** – provides understanding of a model's overall behavior across the entire dataset

# Local Interpretable Model-agnostic Explanations (LIME)

Sometimes you don't know if you can trust a machine learning prediction...



# Local Interpretable Model-agnostic Explanations (LIME)

- LIME creates local explanations
- It follows this general set of steps to create explanations

# Local Interpretable Model-agnostic Explanations (LIME)

- LIME creates local explanations
- It follows this general set of steps to create explanations
  - Pick a specific data point that you want explained

# Local Interpretable Model-agnostic Explanations (LIME)

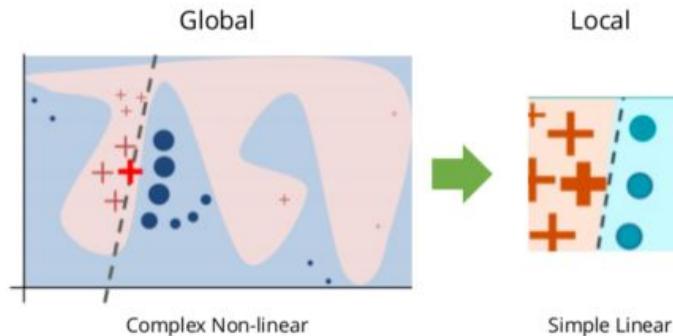
- LIME creates local explanations
- It follows this general set of steps to create explanations
  - Pick a specific data point that you want explained
  - LIME creates a **perturbed** dataset of that sample

# Local Interpretable Model-agnostic Explanations (LIME)

- LIME creates local explanations
- It follows this general set of steps to create explanations
  - Pick a specific data point that you want explained
  - LIME creates a **perturbed** dataset of that sample
  - Your black box model is used to make predictions on the **perturbed** dataset

# Local Interpretable Model-agnostic Explanations (LIME)

- LIME creates local explanations
- It follows this general set of steps to create explanations
  - Pick a specific data point that you want explained
  - LIME creates a **perturbed** dataset of that sample
  - Your black box model is used to make predictions on the **perturbed** dataset
  - LIME trains a surrogate model based on the model predictions

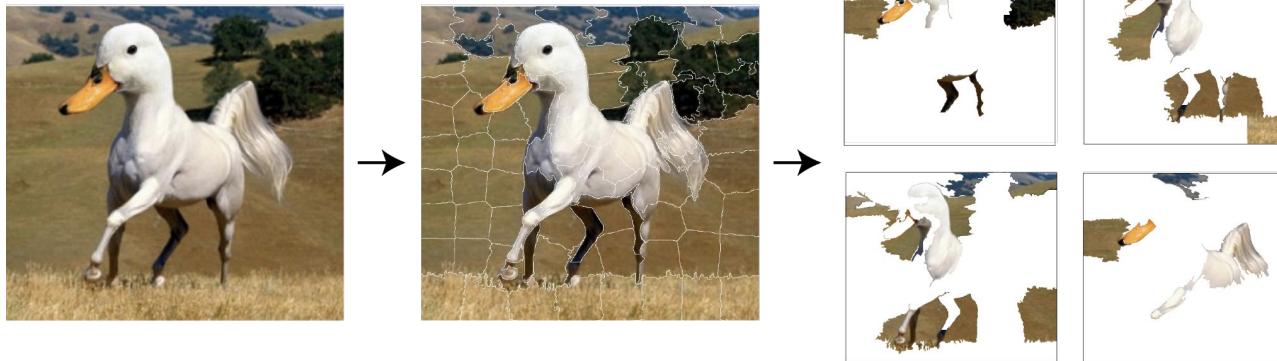


# Local Interpretable Model-agnostic Explanations (LIME)

- LIME creates local explanations
- It follows this general set of steps to create explanations
  - Pick a specific data point that you want explained
  - LIME creates a **perturbed** dataset of that sample
  - Your black box model is used to make predictions on the **perturbed** dataset
  - LIME trains a surrogate model based on the model predictions
  - Surrogate is datamined to determine local feature importance

# Local Interpretable Model-agnostic Explanations (LIME)

- LIME creates local explanations
- It follows this general set of steps to create explanations
  - Pick a specific data point that you want explained
  - LIME creates a **perturbed** dataset of that sample
  - Your black box model is used to make predictions on the **perturbed** dataset
  - LIME trains a surrogate model based on the model predictions
  - Surrogate is datamined to determine local feature importance
  - LIME creates a human readable output



**Label: standard poodle**

**Probability: 0.18**

**Explanation Fit: 0.37**



**Label: goose**

**Probability: 0.15**

**Explanation Fit: 0.55**



# SHapley Additive exPlanations (SHAP)

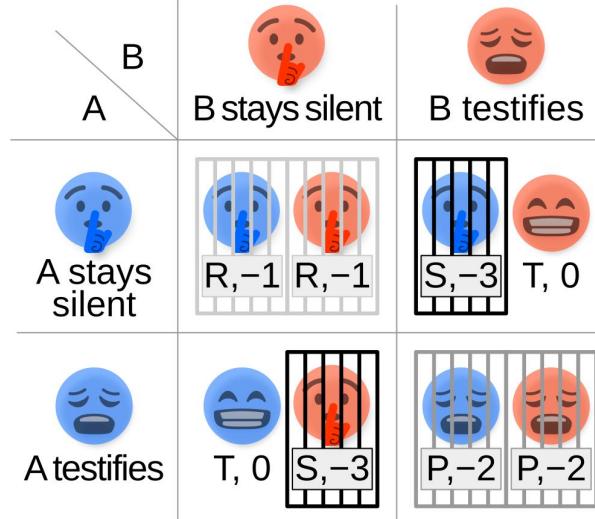
- Based off of cooperative game-theory

# SHapley Additive exPlanations (SHAP)

- Based off of cooperative game-theory
  - **Players** – features in a model
  - **Coalitions** – feature subsets
  - **Payoff** – model predictions for an instance

# SHapley Additive exPlanations (SHAP)

- Based off of cooperative game-theory
  - **Players** – features in a model
  - **Cohorts** – feature subsets
  - **Payoff** – model predictions for an instance
- Famous example of cooperative game-theory
- Prisoner's Dilemma
  - Two rational agents
  - Can either cooperate or stay silent
  - Dilemma is that the payoff is higher if the agents cooperate



# SHapley Additive exPlanations (SHAP)

- SHAP can create both local and global explanations

# SHapley Additive exPlanations (SHAP)

- SHAP can create both local and global explanations
- It follows this general set of steps to create explanations

# SHapley Additive exPlanations (SHAP)

- SHAP can create both local and global explanations
- It follows this general set of steps to create explanations
  - Uses the model to make a baseline prediction (without features)

# SHapley Additive exPlanations (SHAP)

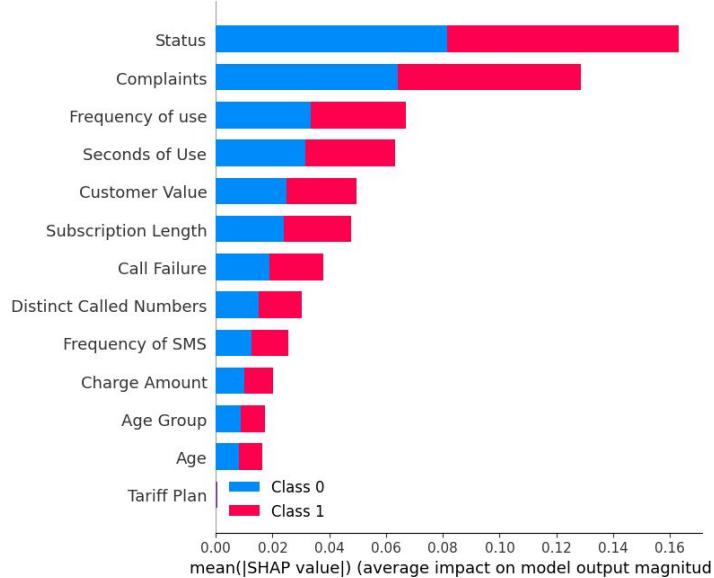
- SHAP can create both local and global explanations
- It follows this general set of steps to create explanations
  - Uses the model to make a baseline prediction (without features)
  - For every feature
    - Consider all subsets (coalitions) that exclude that feature
    - Add that feature to each coalition and calculate the change in model predictions

# SHapley Additive exPlanations (SHAP)

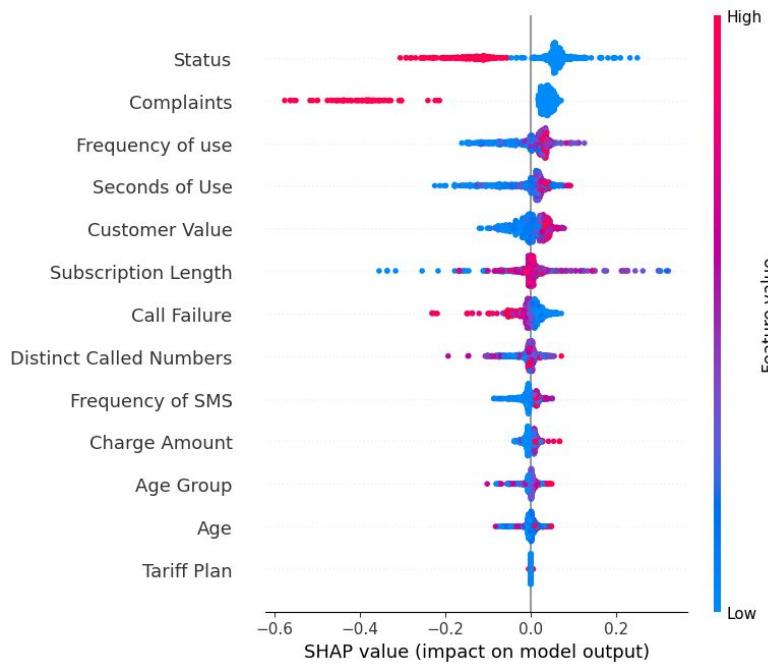
- SHAP can create both local and global explanations
- It follows this general set of steps to create explanations
  - Uses the model to make a baseline prediction (without features)
  - For every feature
    - Consider all subsets (coalitions) that exclude that feature
    - Add that feature to each coalition and calculate the change in model predictions
  - Compute the marginal contributions of the selected feature

# SHapley Additive exPlanations (SHAP)

- SHAP can create both local and global explanations
- It follows this general set of steps to create explanations
  - Uses the model to make a baseline prediction (without features)
  - For every feature
    - Consider all subsets (coalitions) that exclude that feature
    - Add that feature to each coalition and calculate the change in model predictions
  - Compute the marginal contributions of the selected feature
  - Using the marginal contributions, calculate the additive shapley value

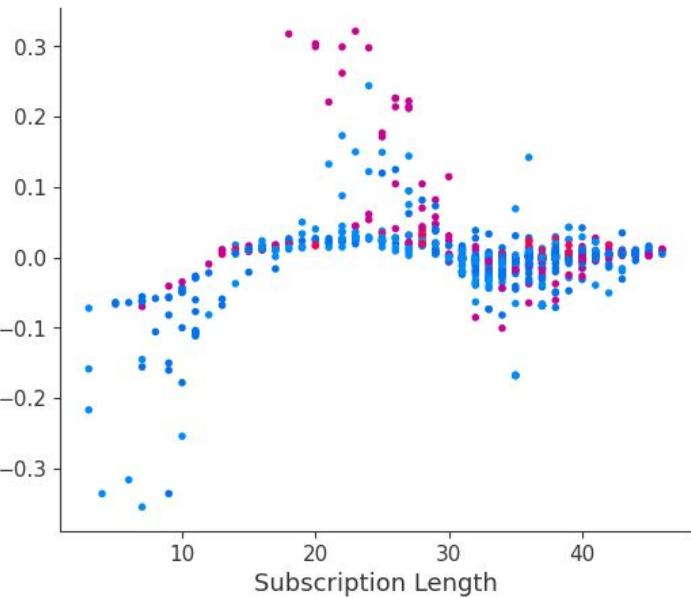


**Summary Plot**

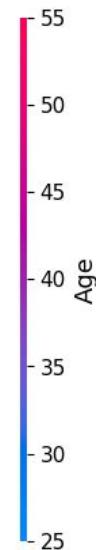


**Other Summary Plot**

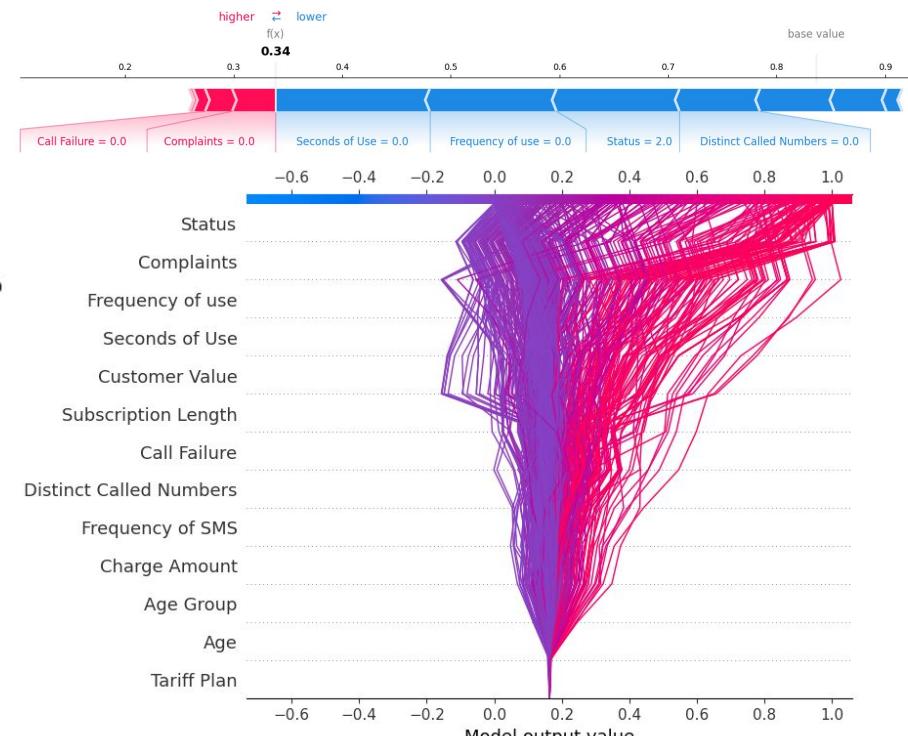
SHAP value for Subscription Length



**Dependence Plot  
For Subscription Length**



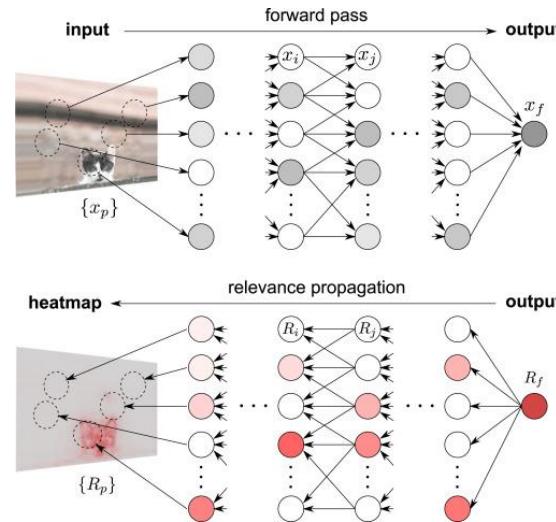
**Force Plot**



**Decision Plot**

# Layer-wise Relevance Propagation (LRP)

- LRP creates local explanations specifically for neural networks
- It follows this general set of steps to create explanations



# Layer-wise Relevance Propagation (LRP)

- LRP creates local explanations specifically for neural networks
- It follows this general set of steps to create explanations
  - Forward pass a sample through the neural network (create a prediction)

# Layer-wise Relevance Propagation (LRP)

- LRP creates local explanations specifically for neural networks
- It follows this general set of steps to create explanations
  - Forward pass a sample through the neural network (create a prediction)
  - Assign a relevance score equal to the prediction value

# Layer-wise Relevance Propagation (LRP)

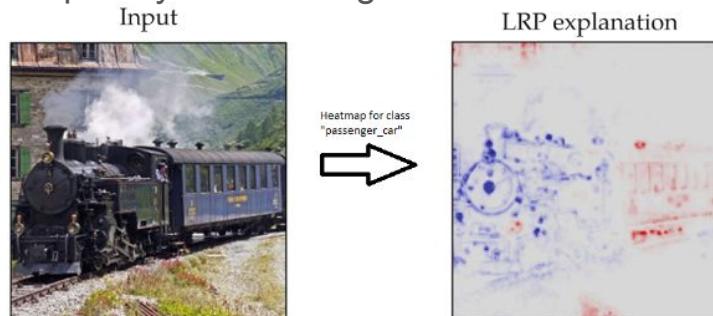
- LRP creates local explanations specifically for neural networks
- It follows this general set of steps to create explanations
  - Forward pass a sample through the neural network (create a prediction)
  - Assign a relevance score equal to the prediction value
  - Relevance score is propagated back through the network
    - Each layer redistributes the relevance it received among its neurons proportional to their contribution

# Layer-wise Relevance Propagation (LRP)

- LRP creates local explanations specifically for neural networks
- It follows this general set of steps to create explanations
  - Forward pass a sample through the neural network (create a prediction)
  - Assign a relevance score equal to the prediction value
  - Relevance score is propagated back through the network
    - Each layer redistributes the relevance it received among its neurons proportional to their contribution
    - There are redistribution rules used to determine how this happens

# Layer-wise Relevance Propagation (LRP)

- LRP creates local explanations specifically for neural networks
- It follows this general set of steps to create explanations
  - Forward pass a sample through the neural network (create a prediction)
  - Assign a relevance score equal to the prediction value
  - Relevance score is propagated back through the network
    - Each layer redistributes the relevance it received among its neurons proportional to their contribution
    - There are redistribution rules used to determine how this happens
  - Propagate back to input layer and assign relevance scores



# Quick Review

- Perturbation based explanations (LIME)
- Game-theory/Feature based explanations (SHAP)
- Decomposition based (LRP)

# Quick Review

- Perturbation based explanations (LIME)
- Game-theory/Feature based explanations (SHAP)
- Decomposition based (LRP)
- Python packages
  - LIME (pip install lime)
  - SHAP (pip install shap)
  - LRP (no simple pip install, there are other implementations it seems)
  - ELI5 (pip install eli5)
  - Interpret (pip install interpret) – appears to be an explanation suite
  - OmniXAI (pip install omnixai) – appears to be an explanation suite

# The Counterfactual

- Local explanation centered around flipping the label
  - “Explanation describes the smallest change to the feature values that changes the prediction to a predefined output”
- Becoming more common in the literature (especially with the rise of LLMs)

# The Counterfactual

- Local explanation centered around flipping the label
  - “Explanation describes the smallest change to the feature values that changes the prediction to a predefined output”
- Becoming more common in the literature (especially with the rise of LLMs)
- General Algorithm:
  - Select a sample to be explained and the desired label to change to ( $y'$ )

# The Counterfactual

- Local explanation centered around flipping the label
  - “Explanation describes the smallest change to the feature values that changes the prediction to a predefined output”
- Becoming more common in the literature (especially with the rise of LLMs)
- General Algorithm:
  - Select a sample to be explained and the desired label to change to ( $y'$ )
  - Sample a random instance as the initial counterfactual

# The Counterfactual

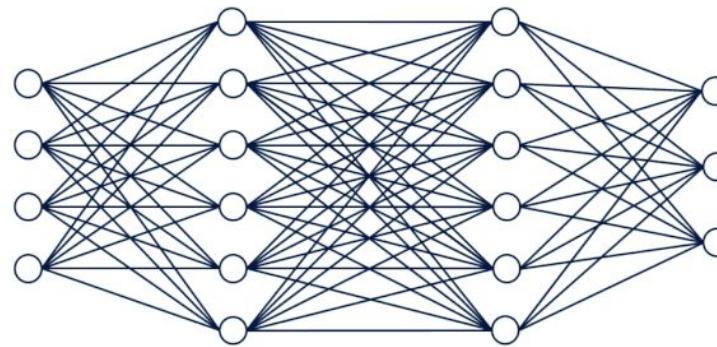
- Local explanation centered around flipping the label
  - “Explanation describes the smallest change to the feature values that changes the prediction to a predefined output”
- Becoming more common in the literature (especially with the rise of LLMs)
- General Algorithm:
  - Select a sample to be explained and the desired label to change to ( $y'$ )
  - Sample a random instance as the initial counterfactual
  - Optimize a loss function with the above counterfactual as the starting point
    - $\rightarrow L(x, x', y', \lambda) = \lambda \cdot (\hat{f}(x') - y')^2 + d(x, x')$

# The Counterfactual

- Local explanation centered around flipping the label
  - “Explanation describes the smallest change to the feature values that changes the prediction to a predefined output”
- Becoming more common in the literature (especially with the rise of LLMs)
- General Algorithm:
  - Select a sample to be explained and the desired label to change to ( $y'$ )
  - Sample a random instance as the initial counterfactual
  - Optimize a loss function with the above counterfactual as the starting point
    - $\rightarrow L(x, x', y', \lambda) = \lambda \cdot (\hat{f}(x') - y')^2 + d(x, x')$
    - $(\hat{f}(x') - y')^2$  – quadratic distance between counterfactual and desired output
    - $d(x, x')$  – distance between original and counterfactual samples
    - $\lambda$  – affects explanations. Higher values prefers predictions close to  $y'$ . Lower values prefer counterfactuals  $x'$  more similar to  $x$

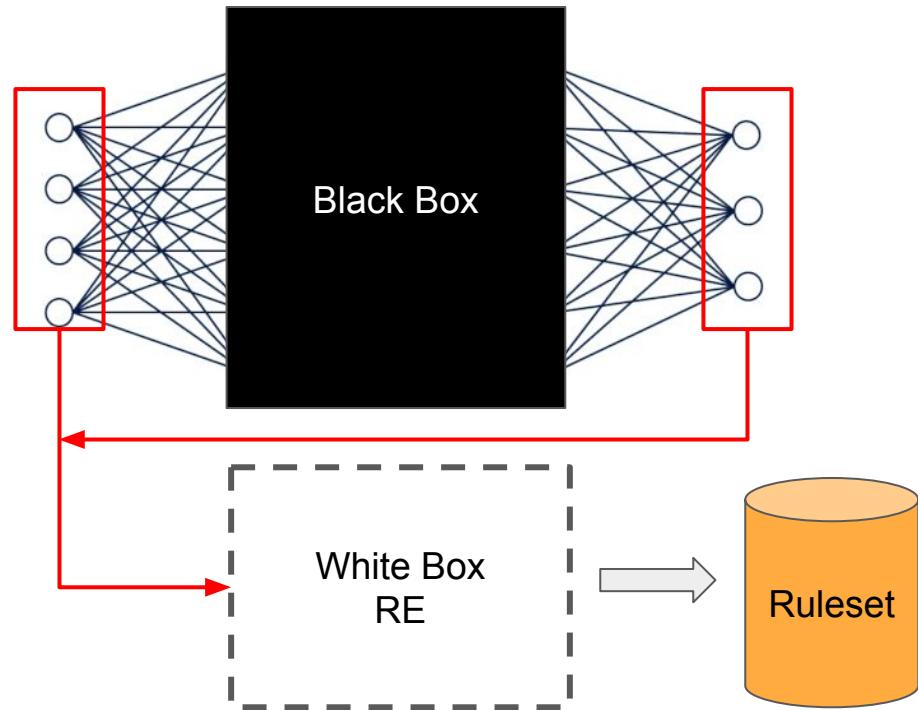
# Rule Extraction

- Types
  - Pedagogical
  - Decompositional
  - Eclectic



# Rule Extraction

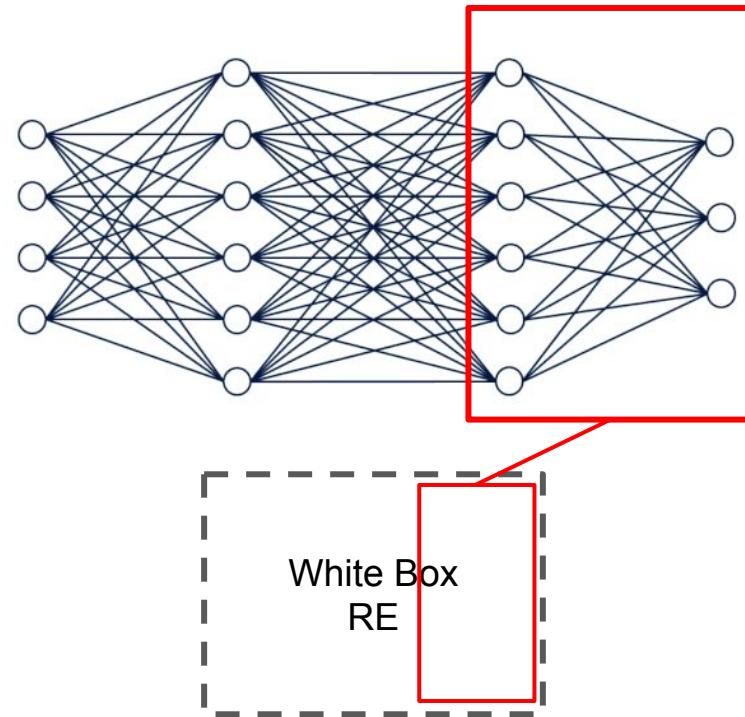
- **Pedagogical**
  - Inputs and outputs
  - Train a decision tree
  - **Maintains the black box**
  - Trustworthy?



# Rule Extraction

- **Decompositional**

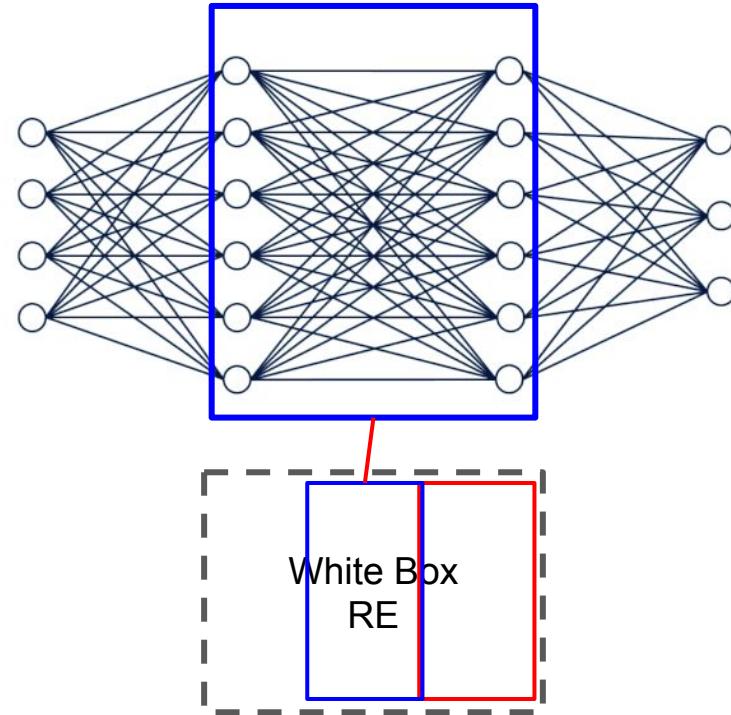
- Uses weights
- Trains a decision tree per layer
- **Opens the black box**
- **Expensive**, but trustworthy



# Rule Extraction

- **Decompositional**

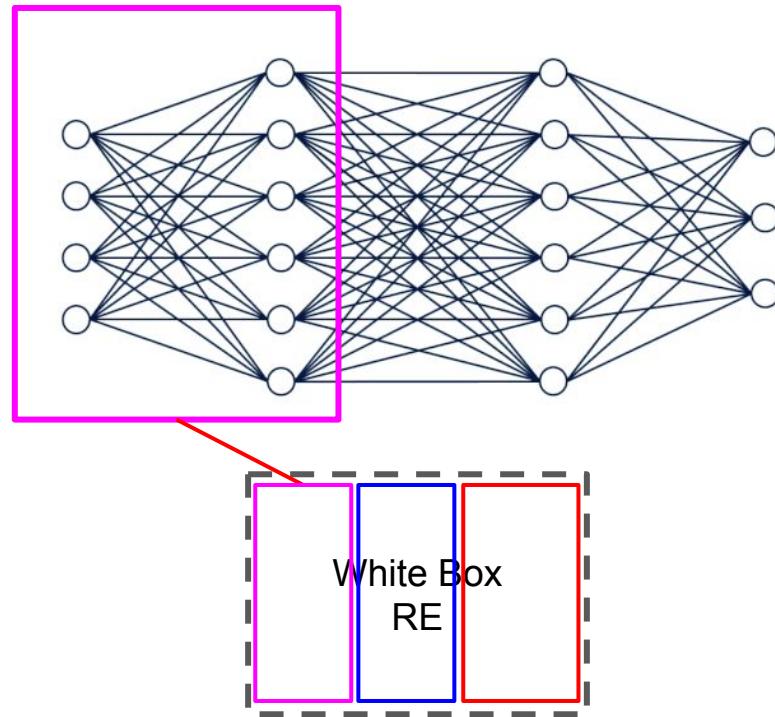
- Uses weights
- Trains a decision tree per layer
- **Opens the black box**
- **Expensive**, but trustworthy



# Rule Extraction

- **Decompositional**

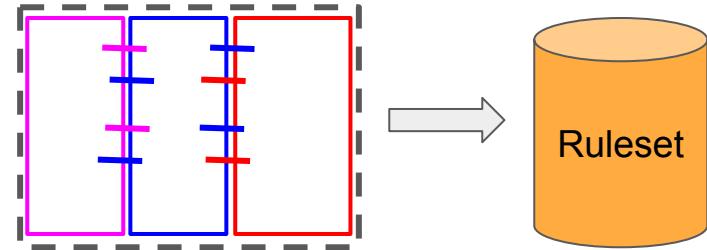
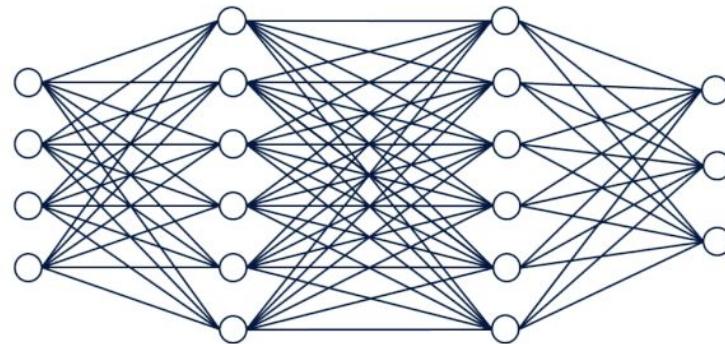
- Uses weights
- Trains a decision tree per layer
- **Opens the black box**
- **Expensive**, but trustworthy



# Rule Extraction

- **Decompositional**

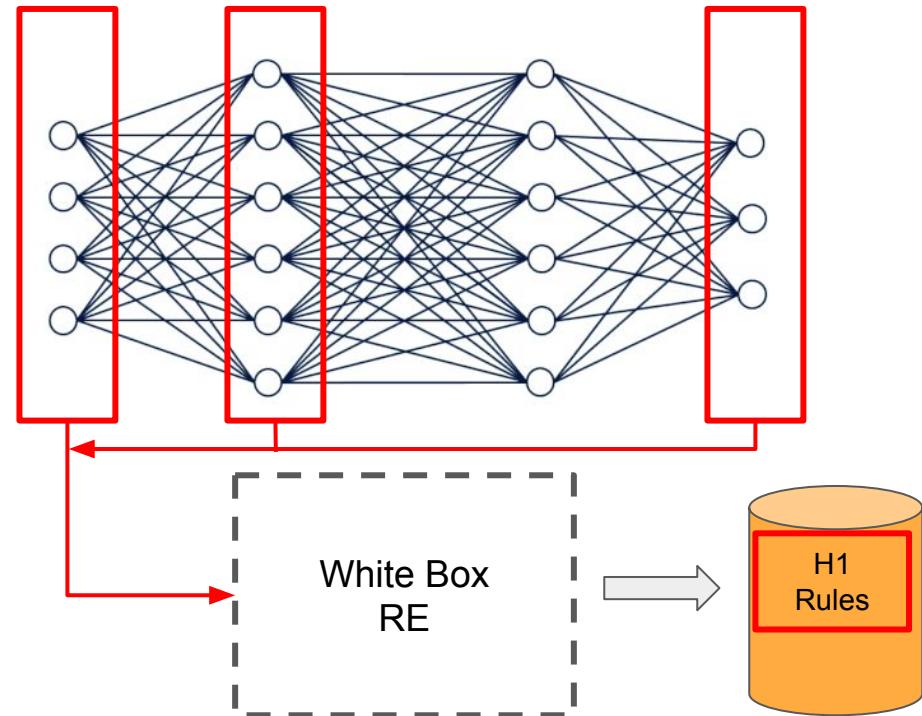
- Uses weights
- Trains a decision tree per layer
- **Opens the black box**
- **Expensive**, but trustworthy



# Rule Extraction

- **Eclectic**

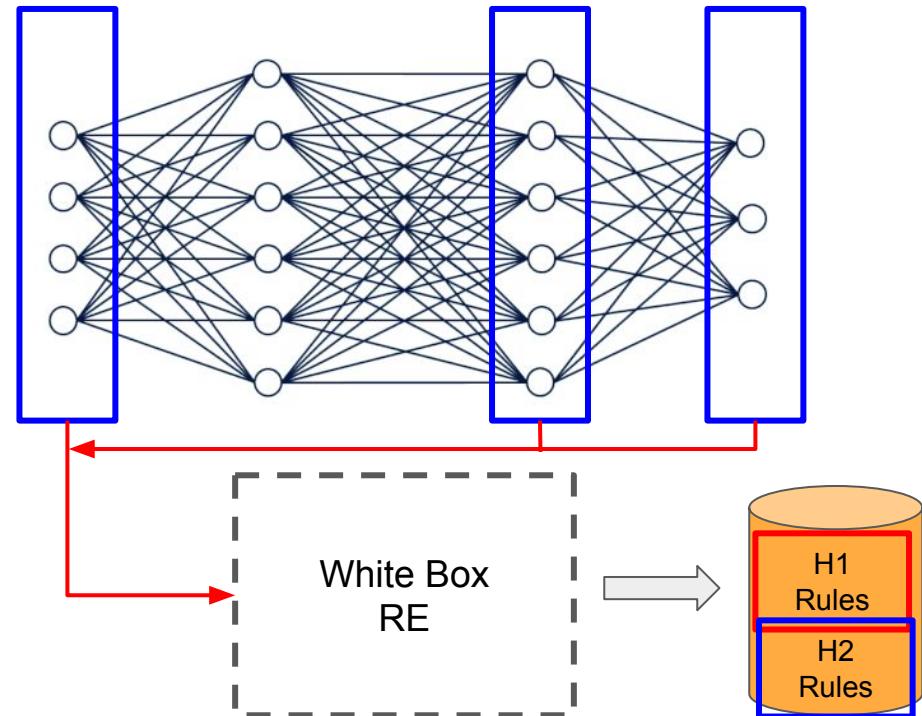
- Uses Inputs, weights, outputs
- Extracts rules per layer
- **Opens the black box**
- **Not as Expensive**, but trustworthy



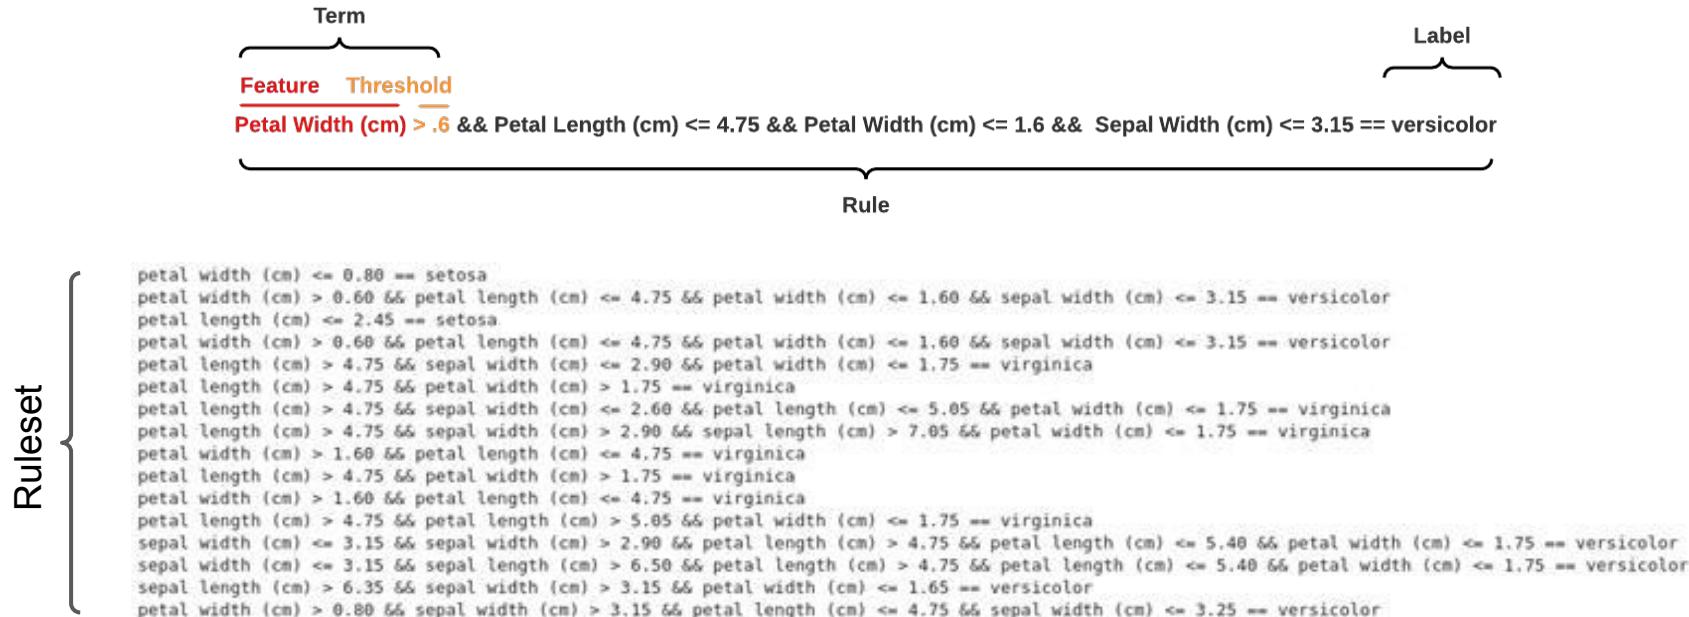
# Rule Extraction

- **Eclectic**

- Uses Inputs, weights, outputs
- Extracts rules per layer
- **Opens the black box**
- **Not as Expensive**, but trustworthy



# Rule Extraction



# Eclectic Rule Extraction Algorithm

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

1:  $R = set()$

2:  $Y' = M.predict(X)$

3: **for** hidden layer  $h_i$  in  $M$  **do**

4:    $X' = h_i(X)$

5:    $hidden_{dt} = DT_{alg}(X', Y')$

6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$

7:    $\hat{Y} = Rules_{hidden}(X')$

8:    $input_{dt} = DT_{alg}(X, \hat{Y})$

9:    $R.add(ExtractRules(input_{dt}))$

10: **end for**

11: **return**  $R$

**END**

---

# Eclectic Rule Extraction Algorithm

- Input a dataset, model, and decision tree algorithm

---

## Algorithm 1 DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

- 1:  $R = set()$
- 2:  $Y' = M.predict(X)$
- 3: **for** hidden layer  $h_i$  in  $M$  **do**
- 4:    $X' = h_i(X)$
- 5:    $hidden_{dt} = DT_{alg}(X', Y')$
- 6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$
- 7:    $\hat{Y} = Rules_{hidden}(X')$
- 8:    $input_{dt} = DT_{alg}(X, \hat{Y})$
- 9:    $R.add(ExtractRules(input_{dt}))$

10: **end for**

11: **return**  $R$

**END**

---

# Eclectic Rule Extraction Algorithm

- Input a dataset, model, and decision tree algorithm
- Initialize a set to save rulesets

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

```
1:  $R = set()$ 
2:  $Y' = M.predict(X)$ 
3: for hidden layer  $h_i$  in  $M$  do
4:    $X' = h_i(X)$ 
5:    $hidden_{dt} = DT_{alg}(X', Y')$ 
6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$ 
7:    $\hat{Y} = Rules_{hidden}(X')$ 
8:    $input_{dt} = DT_{alg}(X, \hat{Y})$ 
9:    $R.add(ExtractRules(input_{dt}))$ 
```

10: **end for**

11: **return**  $R$

**END**

---

# Eclectic Rule Extraction Algorithm

- Input a dataset, model, and decision tree algorithm
- Initialize a set to save rulesets
- Get dataset predictions from model

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

1:  $R = set()$

2:  $Y' = M.predict(X)$

3: **for** hidden layer  $h_i$  in  $M$  **do**

4:    $X' = h_i(X)$

5:    $hidden_{dt} = DT_{alg}(X', Y')$

6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$

7:    $\hat{Y} = Rules_{hidden}(X')$

8:    $input_{dt} = DT_{alg}(X, \hat{Y})$

9:    $R.add(ExtractRules(input_{dt}))$

10: **end for**

11: **return**  $R$

**END**

---

# Eclectic Rule Extraction Algorithm

- For each hidden layer:

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

1:  $R = set()$

2:  $Y' = M.predict(X)$

3: **for** hidden layer  $h_i$  in  $M$  **do**

4:  $X' = h_i(X)$

5:  $hidden_{dt} = DT_{alg}(X', Y')$

6:  $Rules_{hidden} = ExtractRules(hidden_{dt})$

7:  $\hat{Y} = Rules_{hidden}(X')$

8:  $input_{dt} = DT_{alg}(X, \hat{Y})$

9:  $R.add(ExtractRules(input_{dt}))$

10: **end for**

11: **return**  $R$

**END**

---

# Eclectic Rule Extraction Algorithm

- For each hidden layer:
  - Transform samples in  $X$  using hidden neuron values

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

1:  $R = set()$

2:  $Y' = M.predict(X)$

3: **for** hidden layer  $h_i$  in  $M$  **do**

4:    $X' = h_i(X)$

5:    $hidden_{dt} = DT_{alg}(X', Y')$

6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$

7:    $\hat{Y} = Rules_{hidden}(X')$

8:    $input_{dt} = DT_{alg}(X, \hat{Y})$

9:    $R.add(ExtractRules(input_{dt}))$

10: **end for**

11: **return**  $R$

**END**

---

# Eclectic Rule Extraction Algorithm

- For each hidden layer:
  - Transform samples in  $X$  using hidden neuron values
  - Train a hidden layer decision tree

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

```
1:  $R = set()$ 
2:  $Y' = M.predict(X)$ 
3: for hidden layer  $h_i$  in  $M$  do
4:    $X' = h_i(X)$ 
5:    $hidden_{dt} = DT_{alg}(X', Y')$ 
6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$ 
7:    $\hat{Y} = Rules_{hidden}(X')$ 
8:    $input_{dt} = DT_{alg}(X, \hat{Y})$ 
9:    $R.add(ExtractRules(input_{dt}))$ 
```

10: **end for**

11: **return**  $R$

**END**

---

# Eclectic Rule Extraction Algorithm

- For each hidden layer:
  - Transform samples in  $X$  using hidden neuron values
  - Train a hidden layer decision tree
  - Extract rules from the trained hidden layer decision tree

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

```
1:  $R = set()$ 
2:  $Y' = M.predict(X)$ 
3: for hidden layer  $h_i$  in  $M$  do
4:    $X' = h_i(X)$ 
5:    $hidden_{dt} = DT_{alg}(X', Y')$ 
6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$ 
7:    $Y = Rules_{hidden}(X')$ 
8:    $input_{dt} = DT_{alg}(X, Y)$ 
9:    $R.add(ExtractRules(input_{dt}))$ 
10: end for
11: return  $R$ 
END
```

---

# Eclectic Rule Extraction Algorithm

- For each hidden layer:
  - Transform samples in  $X$  using hidden neuron values
  - Train a hidden layer decision tree
  - Extract rules from the trained hidden layer decision tree
  - Predict using the hidden layer ruleset

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

```
1:  $R = set()$ 
2:  $Y' = M.predict(X)$ 
3: for hidden layer  $h_i$  in  $M$  do
4:    $X' = h_i(X)$ 
5:    $hidden_{dt} = DT_{alg}(X', Y')$ 
6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$ 
7:    $\hat{Y} = Rules_{hidden}(X')$ 
8:    $input_{dt} = DT_{alg}(X, Y)$ 
9:    $R.add(ExtractRules(input_{dt}))$ 
10: end for
11: return  $R$ 
END
```

---

# Eclectic Rule Extraction Algorithm

- For each hidden layer:
  - Transform samples in  $X$  using hidden neuron values
  - Train a hidden layer decision tree
  - Extract rules from the trained hidden layer decision tree
  - Predict using the hidden layer ruleset
  - Train input DT using hidden layer predictions and original dataset

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

```
1:  $R = set()$ 
2:  $Y' = M.predict(X)$ 
3: for hidden layer  $h_i$  in  $M$  do
4:    $X' = h_i(X)$ 
5:    $hidden_{dt} = DT_{alg}(X', Y')$ 
6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$ 
7:    $\hat{Y} = Rules_{hidden}(X')$ 
8:    $input_{dt} = DT_{alg}(X, \hat{Y})$ 
9:    $R.add(ExtractRules(input_{dt}))$ 
10: end for
11: return  $R$ 
END
```

---

# Eclectic Rule Extraction Algorithm

- For each hidden layer:
  - Transform samples in  $X$  using hidden neuron values
  - Train a hidden layer decision tree
  - Extract rules from the trained hidden layer decision tree
  - Predict using the hidden layer ruleset
  - Train input DT using hidden layer predictions and original dataset
  - Extract rules from the input DT and add them to the ruleset

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

```
1:  $R = set()$ 
2:  $Y' = M.predict(X)$ 
3: for hidden layer  $h_i$  in  $M$  do
4:    $X' = h_i(X)$ 
5:    $hidden_{dt} = DT_{alg}(X', Y')$ 
6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$ 
7:    $\hat{Y} = Rules_{hidden}(X')$ 
8:    $input_{dt} = DT_{alg}(X, \hat{Y})$ 
9:    $R.add(ExtractRules(input_{dt}))$ 
10: end for
11: return  $R$ 
END
```

---

# Eclectic Rule Extraction Algorithm

- For each hidden layer:
  - Transform samples in  $X$  using hidden neuron values
  - Train a hidden layer decision tree
  - Extract rules from the trained hidden layer decision tree
  - Predict using the hidden layer ruleset
  - Train input DT using hidden layer predictions and original dataset
  - Extract rules from the input DT and add them to the ruleset
- Return the final ruleset

---

**Algorithm 1** DNN Rule Extraction

---

**Input:** Dataset ( $X$ ), Model ( $M$ ), Decision Tree Algorithm ( $DT_{alg}$ )

**Output:** Final Ruleset ( $R$ )

**BEGIN**

```
1:  $R = set()$ 
2:  $Y' = M.predict(X)$ 
3: for hidden layer  $h_i$  in  $M$  do
4:    $X' = h_i(X)$ 
5:    $hidden_{dt} = DT_{alg}(X', Y')$ 
6:    $Rules_{hidden} = ExtractRules(hidden_{dt})$ 
7:    $\hat{Y} = Rules_{hidden}(X')$ 
8:    $input_{dt} = DT_{alg}(X, \hat{Y})$ 
9:    $R.add(ExtractRules(input_{dt}))$ 
10: end for
11: return  $R$ 
```

**END**

---