
Anomaly Detection for Predictive Maintenance

Jason Sonith, Steve Nguyen
University of South Alabama

Why This Matters

- Unplanned machine failures costs hundreds of thousands per hour
- Small companies and plants can't afford unexpected downtime
- But they also lack the resources needed for modern ML solutions:
 - ML specialists
 - GPU Servers
 - Labeled failure data

So the problem becomes:

“How do you bring predictive maintenance to teams who can’t afford it?”

What's Wrong With Current Solutions?

Deep learning models:

- Require GPU's
- Hard to deploy
- Need large labeled datasets

Threshold alarms:

- Too many false alerts
- Don't adapt to changes

Vendor Solutions:

- Expensive
- Cloud-locked
- Security concerns
- Exclusivity

Our Solution


Our approach:

A lightweight, CPU-only,
YAML-driven anomaly detection
pipeline.

This pipeline:

- Runs on regular laptops
- No machine learning expertise required
- Uses simple YAML configs
- Calibrates false alarms per week
- 4 anomaly detection models
- SHAP explainability
- Works across 4 industrial datasets

Pipeline Architecture



Pipeline Stages:

- Data Prep
- Feature Engineering
- Model Training
- Threshold Calibration
- Evaluation & SHAP
- Production Scoring

How it works

Define in YAML

```
dataset_name: ims
paths:
  raw_input_dir: data/raw/ims/
window:
  size: 2048      # Window length
  stride: 1024    # 50% overlap
split:
  train_ratio: 0.60 # First 60% = healthy
  test_ratio: 0.30 # Last 30% = degraded
computed_features:
  - rms            # Root Mean Square
  - peak_to_peak   # Amplitude range
  - kurtosis       # Tail heaviness
```

All settings in one file

```
# Data preparation
python scripts/prep_data.py --config configs/ims.yaml

# Feature engineering
python scripts/make_features.py --config configs/ims.yaml

# Train model
python scripts/train.py --config configs/models/isolation_forest.yaml

# Calibrate threshold
python scripts/threshold.py --target_far 0.2/week
```

Run Pipeline

```
Loading config from configs/ims.yaml
Dataset: ims
Loading 2156 files...
Loaded 44154880 rows
Columns with NaN values before dropping:
Dropped missing values: 44154880 rows remaining
Saving to data/clean/ims/ims_clean.parquet (Parquet)...
Saved 44154880 rows to parquet
Saving to data/clean/ims/ims_clean.csv (CSV)...
Saved 44154880 rows to csv
```

The Datasets

Datasets Used



IMS Bearings – real vibration until failure

CWRU Bearings – seeded bearing faults

AI4I – synthetic manufacturing cycles

NASA C-MAPSS – turbofan engine degradation

These Datasets Covers



Vibration, tabular, and multivariate time-series data.

Modeling Approach

4 Anomaly Detection Models



Isolation Forest

Fast & Robust



Local Outlier Factor

Density-Based



One-Class SVM

Normal boundary



Autoencoder

Reconstruction-based

Why SHAP is Important to Our Model


What is SHAP?

SHAP is an interpretability framework based on Shapley values that quantifies how each feature contributes to a model's prediction, indicating exactly which inputs increased or decreased the output.

- ❑ Explains feature influence: Shows the contribution of each input feature to the prediction.
- ❑ Improves model trust: Ensures the model bases decisions on meaningful physical signals rather than noise.
- ❑ Validates correctness: Confirms the model is learning the right patterns.
- ❑ Supports real-world deployment: Helps detect data leakage, overfitting, or unexpected model behavior before real-world use.
- ❑ Communicates results clearly: SHAP visualizations make model decisions easy to explain to non-technical audiences

Influence of SHAP on our scores

ML models are black boxes, Engineers don't trust mysterious alerts

Alerts without SHAP	Alert with SHAP	Top 10 Most Important Features																						
<ul style="list-style-type: none"><input type="checkbox"/> Bearing 3 failed:<input type="checkbox"/> <u>Why</u> did it fail?<input type="checkbox"/> <u>Where</u> did it fail?<input type="checkbox"/> Do we <u>trust</u> the alert?	<ul style="list-style-type: none"><input type="checkbox"/> Bearing 3 failed:<input type="checkbox"/> <u>Why</u>: Channel 2 kurtosis spiked<input type="checkbox"/> <u>Where</u>: Sensor on channel 2<input type="checkbox"/> <u>Trust</u>: We can verify the logic	 <table border="1"><thead><tr><th>Feature</th><th>Value</th></tr></thead><tbody><tr><td>ch2_crest_factor</td><td>~0.085</td></tr><tr><td>ch3_crest_factor</td><td>~0.085</td></tr><tr><td>ch1_crest_factor</td><td>~0.08</td></tr><tr><td>ch8_crest_factor</td><td>~0.08</td></tr><tr><td>ch5_crest_factor</td><td>~0.075</td></tr><tr><td>ch4_crest_factor</td><td>~0.055</td></tr><tr><td>ch6_crest_factor</td><td>~0.048</td></tr><tr><td>ch7_crest_factor</td><td>~0.042</td></tr><tr><td>ch2_kurtosis</td><td>~0.04</td></tr><tr><td>ch1_kurtosis</td><td>~0.038</td></tr></tbody></table>	Feature	Value	ch2_crest_factor	~0.085	ch3_crest_factor	~0.085	ch1_crest_factor	~0.08	ch8_crest_factor	~0.08	ch5_crest_factor	~0.075	ch4_crest_factor	~0.055	ch6_crest_factor	~0.048	ch7_crest_factor	~0.042	ch2_kurtosis	~0.04	ch1_kurtosis	~0.038
Feature	Value																							
ch2_crest_factor	~0.085																							
ch3_crest_factor	~0.085																							
ch1_crest_factor	~0.08																							
ch8_crest_factor	~0.08																							
ch5_crest_factor	~0.075																							
ch4_crest_factor	~0.055																							
ch6_crest_factor	~0.048																							
ch7_crest_factor	~0.042																							
ch2_kurtosis	~0.04																							
ch1_kurtosis	~0.038																							

Threshold Calibration

THRESHOLD CALIBRATION RESULTS

Model/Dataset	Target FAR	Estimated FAR	Threshold
IMS IForest	1.0/wk	0.989/wk	0.4851
IMS AutoEncoder	0.2/wk	0.200/wk	0.0137
IMS kNN-LOF	0.2/wk	0.200/wk	1.7821
IMS OC-SVM	2.0/wk	2.000/wk	-0.3182
AI4I IForest	0.2/wk	0.202/wk	0.4863
CWRU IForest	0.2/wk	0.212/wk	0.4795
FD001 IForest	0.2/wk	0.289/wk	0.4912
FD002 IForest	0.2/wk	0.216/wk	0.5025
FD003 IForest	0.2/wk	0.217/wk	0.4933
FD004 IForest	0.2/wk	0.221/wk	0.5016



Goal

**Turn continuous scores
into reliable alerts:**
FAR - False Alarm Rate
(alerts per week)



Calibrate

**Compute threshold based on
validation percentiles.**
We aligned with ISA-18.2 alarm standards.




Results

**Calibration Accuracy:
93-100%**

IMS Bearing Run-to-Failure: Multi-Model Degradation Detection

Markers Show First Detection of Degradation



Highlights

Dataset	Results	Challenges
IMS	<ul style="list-style-type: none"><input type="checkbox"/> All models separated normal vs degraded states<input type="checkbox"/> Isolation Forest had the best overall speed + stability	<ul style="list-style-type: none"><input type="checkbox"/> No labels<input type="checkbox"/> Non-stationary vibration data
CWRU	<ul style="list-style-type: none"><input type="checkbox"/> ROC-AUC: 0.942<input type="checkbox"/> PR-AUC: 0.964<input type="checkbox"/> Precision: 1.000	<ul style="list-style-type: none"><input type="checkbox"/> Low recall due to strict FAR<input type="checkbox"/> Seeded faults ≠ gradual failures
Overall-Cross	<ul style="list-style-type: none"><input type="checkbox"/> Thresholds consistent (~0.48 - 0.50)<input type="checkbox"/> FAR stayed ~0.202 - 0.289/week<input type="checkbox"/> Isolation Forest generalizes best	<ul style="list-style-type: none"><input type="checkbox"/> Windowing/Scaling differed per dataset<input type="checkbox"/> Calibration depends on validation size

Insights and Limitations

Aspect	Explanation	Insight	Limitation
Config-Driven Design	YAML control made the pipeline reusable and easy to adapt without code changes.	✓	
Model-Agnostic Framework	The pipeline worked consistently across models because engineering mattered more than algorithm choice.	✓	
False-Alarm Control is Crucial	Industrial systems care more about predictable alert behavior than maximizing recall	✓	
Real-World Drift Not Modeled	Datasets don't capture long-term sensor movement or changing machine behavior.		✗
Limited Validation Sets	Small validation splits make threshold calibration less stable in some scenarios.		✗
Real-Time Monitoring	Pipeline only handles batch scoring and not real-time conditions.		✗
Autoencoder Complexity	Neural model required more compute + tuning compared to classical methods.		✗

Conclusion / Q&A

- ❑ We developed a lightweight, CPU-only anomaly detection pipeline for predictive maintenance on industrial equipment..
- ❑ A config-driven design allows easy reuse across machines, datasets, and sensors with minimal changes.
- ❑ False-alarm calibration enables predictable alert behavior aligned with industrial standards.
- ❑ The pipeline generalizes across four diverse datasets, demonstrating broad applicability.
- ❑ Classical models + strong engineering proved sufficient for reliable performance.
- ❑ Provides a practical foundation for deployable, real-world maintenance monitoring

Questions?
